



Cyber Physical Systems Security 12th Project

Participants:

- | | |
|-----------------------------------|------------------|
| - Mohamed Mostafa Sayed Saber Ali | 221027679 |
| - Mohamed Hany El Kordy | 221027552 |
| - Basel Hany Abbas | 221027635 |
| - Mohamed Basem | 221027617 |
| - Mohamed Mostafa Abd El Baset | 221027785 |
| - Omar Hesham | 221027739 |
-

Table of Contents:

- **Project Objectives**
 - o Understand security threats
 - o Study AI-based IDS techniques
 - o Design and evaluate IDS
- **Literature Review**
 - o IDS in IoT networks
 - o Machine learning & deep learning for IDS
 - o Common IoT attacks (DoS, MITM, etc.)
 - o Public IDS datasets
- **Task Description / Implementation**
 - o IoT system architecture & attack points
 - o Analysis of Cyber-Attacks
 - o AI-Based IDS Design & Detection Workflow
 - o Evaluation metrics
 - o Advantages and Limitations
- **Results / Simulation**
 - o Sample detection results
- **Conclusion**
 - o Summary of results
 - o Future improvements

Project Objectives

- The main goal of this project is to **design** and **implement** an **AI-based Intrusion Detection System (IDS)** for **IoT networks**. IoT devices and Cyber-Physical Systems (CPS) face unique **security challenges** due to their **limited computational resources**, **heterogeneous devices**, and **large attack surface**. Traditional rule-based security mechanisms often fail to detect sophisticated or evolving attacks. Therefore, this project focuses on developing an intelligent and adaptive IDS using machine learning and deep learning techniques.



- The specific objectives of this project are:
 - o **Understand Security Threats in IoT-based CPS Environments**
 - Identify common vulnerabilities in IoT devices and networks.
 - Analyze the types of cyber-attacks that IoT systems are susceptible to, such as Denial-of-Service (DoS), Man-in-the-Middle (MITM), Eavesdropping, and Data Injection Attacks.
 - o **Study AI-based Intrusion Detection Techniques**
 - Explore machine learning and deep learning approaches for detecting anomalies and attacks in network traffic.

- Evaluate the suitability of different algorithms for real-time IoT network security.
- **Design and Evaluate an AI-based IDS Model**
 - Develop a **deep learning model** that can classify **network traffic** as **normal** or **malicious**.
 - **Preprocess** and **scale IoT network datasets** for **model training**.
 - Evaluate the **effectiveness** of the **model** using **metrics** such as **accuracy**, **precision**, **recall**, and **F1-score**.
 - Analyze the **performance**, **limitations**, and **potential improvements** for **real-world deployment**.


Literature Review

- **Intrusion Detection Systems (IDS)** play a critical role in **securing IoT networks**, which are increasingly targeted by **cyber-attacks** due to their widespread deployment and limited device resources. This section summarizes existing research and approaches relevant to AI-based IDS for IoT networks.



- IDS in IoT Networks:


- o IoT networks consist of **heterogeneous devices** communicating over **various protocols**, which introduces **unique security challenges**.
- o Traditional IDS approaches, such as **signature-based** or **rule-based systems**, often fail to detect **new** or **evolving attacks**.
- o **AI-based IDS** leverage intelligent algorithms to analyze network traffic and detect anomalies or malicious behavior in real-time.



- Machine Learning and Deep Learning for IDS:

- o **Machine Learning (ML) techniques**, such as **Decision Trees**, **Random Forests**, and **Support Vector Machines**, can classify network traffic based on statistical patterns.
- o **Deep Learning (DL) approaches**, such as **Artificial Neural Networks (ANNs)**, **Convolutional Neural Networks (CNNs)**, and **Recurrent Neural Networks (RNNs)**, can automatically extract features from complex data and achieve higher accuracy in detecting unknown attacks.

- AI-based IDS models can adapt to changing network conditions and evolving attack patterns, which is essential for IoT security.



- Common IoT Attacks:

- IoT networks are **susceptible to several types of cyber-attacks:**
 - **Denial-of-Service (DoS):** Flooding the network or devices with excessive requests, causing them to crash or become unresponsive.
 - **Man-in-the-Middle (MITM):** Intercepting and potentially altering communications between IoT devices.
 - **Eavesdropping:** Unauthorized monitoring of data transmissions.
 - **Data Injection Attacks:** Injecting malicious or false data into IoT systems to manipulate or disrupt operations.
 - **Malware-based attacks (e.g., Mirai):** Infecting IoT devices to create botnets or launch coordinated attacks.

OWASP **TOP 10** INTERNET OF THINGS 2018

- 1 Weak, Guessable, or Hardcoded Passwords**
Use of easily bruteforced, publicly available, or unchangeable credentials, including backdoors in firmware or client software that grants unauthorized access to deployed systems.

- 2 Insecure Network Services**
Unneeded or insecure network services running on the device itself, especially those exposed to the internet, that compromise the confidentiality, integrity/authenticity, or availability of information or allow unauthorized remote control...

- 3 Insecure Ecosystem Interfaces**
Insecure web, backend API, cloud, or mobile interfaces in the ecosystem outside of the device that allows compromise of the device or its related components. Common issues include a lack of authentication/authorization, lacking or weak encryption, and a lack of input and output filtering.

- 4 Lack of Secure Update Mechanism**
Lack of ability to securely update the device. This includes lack of firmware validation on device, lack of secure delivery (un-encrypted in transit), lack of anti-rollback mechanisms, and lack of notifications of security changes due to updates.

- 5 Use of Insecure or Outdated Components**
Use of deprecated or insecure software components/libraries that could allow the device to be compromised. This includes insecure customization of operating system platforms, and the use of third-party software or hardware components from a compromised supply chain.

- 6 Insufficient Privacy Protection**
User's personal information stored on the device or in the ecosystem that is used insecurely, improperly, or without permission.

- 7 Insecure Data Transfer and Storage**
Lack of encryption or access control of sensitive data anywhere within the ecosystem, including at rest, in transit, or during processing.


- 8 Lack of Device Management**
Lack of security support on devices deployed in production, including asset management, update management, secure decommissioning, systems monitoring, and response capabilities.

- 9 Insecure Default Settings**
Devices or systems shipped with insecure default settings or lack the ability to make the system more secure by restricting operators from modifying configurations.

- 10 Lack of Physical Hardening**
Lack of physical hardening measures, allowing potential attackers to gain sensitive information that can help in a future remote attack or take local control of the device.





- **Mirai Botnet:** Mirai is one of the most infamous IoT malware strains. It infects IoT devices by exploiting default or weak passwords, turning them into a botnet. Once infected, these devices can be used to launch large-scale Distributed Denial-of-Service (DDoS) attacks, disrupting websites, networks, and online services. Variants of Mirai, such as Mirai-Ackflooding or Mirai-Hostbruteforce, target different IoT protocols and services, making it a persistent threat in IoT security.



- Public IDS Datasets:

- For this project, we used the IoTID20 dataset from Kaggle. This dataset contains both normal and attack traffic in IoT networks, including attacks such as DoS, MITM, and scanning. It is publicly available and can be accessed, just press [here](#).




- Dataset details:

- File:** IoT Network Intrusion Dataset.csv (308.53 MB)
- Usability score:** 2.94
- License:** Unknown
- Update frequency:** Not specified

- This dataset allowed us to train and evaluate the AI-based IDS for detecting malicious activity in IoT environments.


Task Description / Implementation

- The IoT system modeled in this project consists of a variety of connected devices such as sensors, cameras, smart appliances, and gateways. These devices communicate with a central network gateway or cloud service. The architecture can be represented as:
 - o **IoT Devices:** Generate network traffic and collect data.
 - o **Edge Gateway / Router:** Aggregates device traffic and forwards it to cloud or monitoring systems.
 - o **Cloud/Server:** Performs data storage, analytics, and intrusion detection.



- **Potential Attack Points:**

- o **Device Interfaces:** Weak authentication and default credentials make devices vulnerable to malware infection, including botnets like **Mirai**.
- o **Communication Channels:** Traffic between devices and gateways may be intercepted or tampered with, allowing **Man-in-the-Middle (MITM) attacks**.
- o **Network Layer:** Overloading devices or gateways can trigger **Denial-of-Service (DoS) attacks**.
- o **Cloud or Server:** Compromised cloud servers can manipulate or inject malicious data into the IoT network.



- Identifying these attack points is critical to designing a robust Intrusion Detection System (IDS) that can monitor and classify abnormal traffic patterns.

- Analysis of Cyber-Attacks:

- This project focuses on detecting multiple attack types commonly observed in IoT networks.
- The **IoTID20 dataset** used contains labeled examples for the following **Attack Categories (Cat Column)** and **Attack Subcategories (Sub-Cat Column)** respectively:

Attack Categories (Cat column)

- **Mirai** – Botnet malware targeting IoT devices
- **DoS** – Denial-of-Service attacks
- **Scan** – Port and network scanning attempts
- **MITM ARP Spoofing** – Man-in-the-Middle attacks using ARP spoofing
- **Normal** – Legitimate traffic

Attack Subcategories (Sub-Cat column)

- **Mirai-Ackflooding** – Floods the network with ACK packets
- **Mirai-Hostbruteforce** – Attempts to brute-force IoT devices
- **Mirai-UDP Flooding** – Sends high-volume UDP packets to overwhelm devices
- **Mirai-HTTP Flooding** – Floods web services on IoT devices
- **DoS-Synflooding** – TCP SYN flood attack
- **Scan Port OS** – Attempts to detect open ports and operating systems
- **Scan Hostport** – Host-level scanning of ports
- **MITM ARP Spoofing** – Intercepts network traffic
- **Normal** – Legitimate network activity

- The **IDS** is trained to classify traffic into **normal** vs. **anomalous**, while also identifying **attack categories** and **subcategories**.

- AI-Based IDS Design & Detection Workflow:

- o The **IDS** was implemented using **deep learning**, leveraging a **neural network** built with **TensorFlow/Keras**. The workflow includes:
 - **Data Preprocessing:** Before feeding the data into the **neural network**, **irrelevant columns** such as **IP addresses**, **timestamps**, and **categorical labels** that do not contribute to detection were **removed**. **Missing values** and **infinite values** were handled to ensure **data integrity**. Additionally, the **target labels** were **encoded numerically**, with **Normal mapped to 1** and **Anomaly mapped to 0**, enabling the network to perform **binary classification**.

```
=====
# DROP NON-USEFUL COLUMNS
=====
DROP_COLS = [
    "Flow_ID", "Src_IP", "Dst_IP", "Timestamp",
    "cat", "Sub_Cat"
]

for col in DROP_COLS:
    if col in df.columns:
        df.drop([col], axis=1, inplace=True)

=====
# CLEAN DATA
=====
df.replace([np.inf, -np.inf], np.nan, inplace=True)
df.dropna(inplace=True)

=====
# LABEL ENCODING
=====
label_encoder = LabelEncoder()
df["Label"] = label_encoder.fit_transform(df["Label"])

# Check the mapping
print("Label mapping:")
for original, encoded in zip(label_encoder.classes_, range(len(label_encoder.classes_))):
    print(f" {original} -> {encoded}")

x = df.drop("Label", axis=1)
y = df["Label"]

# Print the mapping
for original, encoded in zip(label_encoder.classes_, range(len(label_encoder.classes_))):
    print(f" {original} -> {encoded} ({label_encoder.inverse_transform([original])[0]})")
print(f" {label_encoder.classes_.size} total classes")
```

- **Feature Scaling:** To prepare the data for **neural network training**, all **numerical features** were **normalized** using **StandardScaler**. This **standardization** ensures that **each feature has a similar magnitude**, which helps the **network converge faster** and **improves overall performance**.

```
# =====
# SCALING
# =====
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# SAVE SCALER AND LABEL ENCODER
joblib.dump(scaler, "scaler.save")
joblib.dump(label_encoder, "label_encoder.save") # Save label encoder too!
```

- **Neural Network Architecture:** The IDS model was implemented as a **deep neural network** using **TensorFlow/Keras**. The input layer accepts **all preprocessed features**. **Two hidden layers** are employed: the **first contains 128 neurons** with **ReLU activation** and a **dropout rate of 0.3** to reduce **overfitting**, while the **second has 64 neurons** with **ReLU activation** and the **same dropout rate**. The **output layer** consists of a **single neuron** with a **Sigmoid activation function**, producing a **probability** indicating whether a **network sample is normal or represents an attack**.

```
=====
# IDS MODEL (DEEP LEARNING)
=====

model = Sequential([
    Dense(128, activation="relu", input_shape=(X_train.shape[1],)),
    Dropout(0.3),
    Dense(64, activation="relu"),
    Dropout(0.3),
    Dense(1, activation="sigmoid")
])

model.compile(
    optimizer="adam",
    loss="binary_crossentropy",
    metrics=["accuracy"]
)

early_stop = EarlyStopping(
    monitor="val_loss",
    patience=3,
    restore_best_weights=True
)

model.fit(
    X_train, y_train,
    epochs=15,
    batch_size=1024,
    validation_split=0.2,
    callbacks=[early_stop],
    verbose=1
)

=====
```

- **Training Process:** The dataset was split into **training** and **testing sets**, with **80%** of the data used for **training** and **20%** reserved for **testing**. The network was trained using the **Adam** optimizer with **binary cross-entropy** as the loss function. **Early stopping** was applied to prevent overfitting by monitoring the validation loss and restoring the best model weights. **Training** was conducted over **15 epochs** with a **batch size** of **1024** to balance **convergence speed** and **stability**.

```
# =====
# TRAIN / TEST SPLIT
# =====
x_train, x_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,
    random_state=42,
    stratify=y
)
model.compile(
    optimizer="adam",
    loss="binary_crossentropy",
    metrics=["accuracy"]
)
early_stop = EarlyStopping(
    monitor="val_loss",
    patience=3,
    restore_best_weights=True
)
model.fit(
    x_train, y_train,
    epochs=15,
    batch_size=1024,
    validation_split=0.2,
    callbacks=[early_stop],
    verbose=1
)
#
```

- **Detection Workflow:** During inference, **new network traffic samples are preprocessed and scaled using the same procedures** as for the training data. The **neural network predicts the probability of a sample being normal or malicious**. A **threshold of 0.5** is used for **classification**: **probabilities above 0.5 indicate normal traffic**, while **probabilities below 0.5 indicate an attack**. When an attack is detected, the system also **outputs the corresponding attack category and subcategory**, enabling detailed analysis of the type of intrusion.

```
# =====
# IDS DETECTION FUNCTION
# =====
def detect_intrusion(sample, cat, sub_cat):
    sample = np.array(sample).reshape(1, -1)
    sample = scaler.transform(sample)
    prediction = model.predict(sample, verbose=0)[0][0]

    # Debug: print prediction probability
    print(f"Prediction: {prediction:.4f} (1=Normal, 0=Anomaly)", end=" - ")

    if prediction > 0.5:
        return f"NORMAL TRAFFIC"
    else:
        return f"ATTACK DETECTED [{cat}] ({sub_cat})"
```

- **Evaluation Metrics:**

- o To evaluate IDS performance, the following metrics were used:
 - **Accuracy:** Measures the proportion of correct predictions out of all samples.
 - **Precision:** Fraction of correctly predicted attacks out of all predicted attacks.
 - **Recall (Detection Rate):** Fraction of correctly detected attacks out of all actual attacks.
 - **F1-Score:** Harmonic mean of precision and recall, balancing false positives and false negatives.
 - **Confusion Matrix:** Shows the distribution of true vs. predicted labels:

| | | Actual | |
|-----------|---|----------------------|----------------------|
| | | 1 | 0 |
| | | True Positives (TP) | False Positives (FP) |
| Predicted | 1 | True Positives (TP) | False Positives (FP) |
| | 0 | False Negatives (FN) | True Negatives (TN) |

| | | Actual | |
|-----------|---|----------------------|----------------------|
| | | 1 | 0 |
| | | True Positives (TP) | False Positives (FP) |
| Predicted | 1 | True Positives (TP) | False Positives (FP) |
| | 0 | False Negatives (FN) | True Negatives (TN) |

- **Results:**
 - Achieved ~99.6% accuracy on the test set
 - Minimal false positives (normal traffic misclassified as attack)

```
Accuracy: 0.996338431281629
      precision    recall   f1-score   support
Anomaly          1.00     1.00     1.00    117068
Normal           1.00     0.94     0.97     8015
accuracy          -         -       1.00    125083
macro avg        1.00     0.97     0.98    125083
weighted avg     1.00     1.00     1.00    125083

Model, scaler, and label encoder saved.
```

Model, scaler, and label encoder saved.

- Advantages and Limitations:

- **Advantages:**
 - High accuracy in detecting attacks in IoT networks
 - Scalable to large-scale IoT environments
- **Limitations:**
 - Requires labeled datasets for supervised learning
 - May not detect extremely rare or new attack types
 - Training can be resource-intensive for large networks

Results / Simulation

- To evaluate the **performance** of the **AI-based IDS**, we tested the **trained model** on a mixture of **normal** and **attack traffic samples** from the **IoTID20 dataset**. For each sample, the **model predicts** the **probability** of the **traffic being normal** or an attack. Using a **classification threshold** of **0.5**, the model correctly identified all **types of attacks**, including **Mirai-based attacks**, **Denial-of-Service (DoS) attempts**, and **network scans**, as well as **normal traffic**.
- For example, **normal traffic samples** were consistently classified as **NORMAL TRAFFIC**, while **attack samples** were correctly labeled with their corresponding **category** and **subcategory**, such as **ATTACK DETECTED → Mirai (Mirai-Ackflooding)** or **ATTACK DETECTED →**

DoS (DoS-Synflooding). The predictions also include the **confidence level**, allowing network administrators to assess the **certainty of each detection**.

```
=====
Testing IDS on mixed samples:
=====

Testing 5 Normal samples and 5 Attack samples:

Prediction: 0.9635 (1=Normal, 0=Anomaly) - Sample 1: NORMAL TRAFFIC → [True: Normal] ✓ CORRECT
Prediction: 0.9962 (1=Normal, 0=Anomaly) - Sample 2: NORMAL TRAFFIC → [True: Normal] ✓ CORRECT
Prediction: 0.9998 (1=Normal, 0=Anomaly) - Sample 3: NORMAL TRAFFIC → [True: Normal] ✓ CORRECT
Prediction: 1.0000 (1=Normal, 0=Anomaly) - Sample 4: NORMAL TRAFFIC → [True: Normal] ✓ CORRECT
Prediction: 0.8945 (1=Normal, 0=Anomaly) - Sample 5: NORMAL TRAFFIC → [True: Normal] ✓ CORRECT
Prediction: 0.0000 (1=Normal, 0=Anomaly) - Sample 6: ATTACK DETECTED → Mirai (Mirai-Ackflooding) [True: Attack] ✓ CORRECT
Prediction: 0.0000 (1=Normal, 0=Anomaly) - Sample 7: ATTACK DETECTED → DoS (DoS-Synflooding) [True: Attack] ✓ CORRECT
Prediction: 0.0012 (1=Normal, 0=Anomaly) - Sample 8: ATTACK DETECTED → Scan (Scan Port OS) [True: Attack] ✓ CORRECT
Prediction: 0.0006 (1=Normal, 0=Anomaly) - Sample 9: ATTACK DETECTED → Mirai (Mirai-Hostbruteforceg) [True: Attack] ✓ CORRECT
Prediction: 0.0008 (1=Normal, 0=Anomaly) - Sample 10: ATTACK DETECTED → Mirai (Mirai-Hostbruteforceg) [True: Attack] ✓ CORRECT

=====
Accuracy: 10/10 = 100.0%
=====
```

- A **bulk testing evaluation** on a **larger set of 1,000 randomly selected samples** demonstrated **high accuracy**, with **very few false positives**. The **confusion matrix highlighted** that most **normal and attack samples** were **correctly classified**, and only a minimal number of **normal samples** were **mistakenly flagged as attacks**. This confirms the **effectiveness** of the **IDS** in accurately detecting various **cyber-attacks in IoT networks**.

```
Bulk Testing Accuracy:
=====
Accuracy on 1000 random samples: 0.9980

Confusion Matrix:
      Predicted
      Anomaly  Normal
Actual Anomaly    951     0
Actual Normal      2    47

      precision   recall   f1-score   support
Anomaly        1.00     1.00     1.00     951
Normal         1.00     0.96     0.98     49

      accuracy          1.00     1000
      macro avg       1.00     0.98     0.99     1000
      weighted avg    1.00     1.00     1.00     1000

WARNING: 2 normal samples are being classified as attacks!
```

Conclusion

- The **AI-based IDS** demonstrated **high accuracy** in **detecting both normal and malicious IoT network traffic**. **Sample testing** and **bulk evaluation** confirmed its ability to **correctly classify various attack types**, including **Mirai botnet attacks**, **DoS**, **Scan**, and **MITM attacks**. The results validate the **effectiveness** of using **deep learning** for **intrusion detection in IoT networks**.
- **Future Improvements:**
 - o **Implement real-time IDS deployment**
 - o **Integrate unsupervised learning to detect unknown attacks**



Thank You